

# YOLO (You Only Look Once) Object Detection Tutorial

## Introduction

YOLO (You Only Look Once) is a real-time object detection algorithm that is widely used for detecting objects in images and videos. In this tutorial, we will cover:

- Installing YOLO and dependencies
  - Running YOLO on images and videos
  - Understanding the YOLO model structure
  - Customizing YOLO for your own dataset
- 

## 1. Installing Dependencies

Before using YOLO, ensure you have the necessary dependencies installed.

### Install OpenCV and YOLO Requirements

```
pip install opencv-python numpy torch torchvision
```

### Download Pre-trained YOLO Model

YOLO models can be downloaded from the official sources:

```
wget https://github.com/ultralytics/yolov5/releases/download/v6.0/yolov5s.pt
```

Alternatively, clone the YOLO repository:

```
git clone https://github.com/ultralytics/yolov5.git
cd yolov5
pip install -r requirements.txt
```

---

## 2. Running YOLO on Images

Once YOLO is installed, you can run object detection on an image.

```
python detect.py --weights yolov5s.pt --img 640 --conf 0.25 --source
image.jpg
```

- `--weights yolov5s.pt`: Use the small YOLOv5 model.
- `--img 640`: Set image size to 640 pixels.
- `--conf 0.25`: Confidence threshold for detection.
- `--source image.jpg`: Specify the image file.

The output will display detected objects and save the results in the `runs/detect/` directory.

---

### 3. Running YOLO on Videos or Webcam

You can also use YOLO for real-time detection with a video file or webcam.

#### Detect Objects in a Video File

```
python detect.py --weights yolov5s.pt --source video.mp4
```

#### Run YOLO on a Webcam

```
python detect.py --weights yolov5s.pt --source 0
```

- `--source 0` uses the default webcam.
- 

### 4. Understanding YOLO Model Structure

YOLO uses a neural network to divide an image into grids, predicting bounding boxes and class probabilities for objects in each grid.

- **Backbone:** Extracts features (e.g., CSPDarknet in YOLOv5).
  - **Neck:** Enhances feature extraction (e.g., PANet in YOLOv5).
  - **Head:** Predicts bounding boxes, objectness scores, and class probabilities.
- 

### 5. Custom Training with YOLO

To train YOLO on your own dataset:

#### Step 1: Prepare the Dataset

Your dataset should be in the YOLO format:

```
/images
- train
- val
/labels
- train
- val
```

Each label file should contain:

```
<class_id> <x_center> <y_center> <width> <height>
```

All values are normalized (0-1 range).

## Step 2: Modify the Data Config

Create a `.yaml` file (e.g., `custom_data.yaml`):

```
train: /path/to/train/images
val: /path/to/val/images
nc: 2
names: ['class1', 'class2']
```

- `nc`: Number of classes.
- `names`: List of class names.

## Step 3: Train the Model

```
python train.py --img 640 --batch 16 --epochs 50 --data custom_data.yaml --weights yolov5s.pt
```

- `--img 640`: Image size.
- `--batch 16`: Batch size.
- `--epochs 50`: Number of training epochs.
- `--data custom_data.yaml`: Path to dataset config.
- `--weights yolov5s.pt`: Start training from a pre-trained model.

## Step 4: Test the Model

After training, test your model:

```
python detect.py --weights runs/train/exp/weights/best.pt --source test.jpg
```

---

# Conclusion

This tutorial covered:

- Setting up YOLOv5
- Running YOLO on images and videos
- Understanding YOLO's architecture
- Custom training with YOLO

With this knowledge, you can now customize YOLO for your own object detection tasks. Experiment with different configurations to improve performance!

